



AI reinvention made real

Accenture Innovation Challenge 2026

Team details

TEAM NAME: Accen-Sure



Nishika Haware (Team Leader)

College: IIT (ISM) Dhanbad
Stream: Electronics and Communication Engineering
Year of graduation: 2028



Himanshi Srivastava

College: IIT (ISM) Dhanbad
Stream: Electronics and Communication Engineering
Year of graduation: 2028



Komal Baraskar

College: IIT (ISM) Dhanbad
Stream: Mechanical Engineering
Year of graduation: 2028

PROBLEM STATEMENT:

Every AI deployment carries the same risk: it can be confidently wrong, quietly expensive, or subtly biased — usually discovered only after a user has already acted on it. ControlPlane.ai is a technology layer that continuously observes every AI response in real time across three dimensions: performance, cost, and responsibility. Built to sit on top of any model, it turns AI oversight from an after-the-fact discovery into something you can watch, catch, and act on live.

1

Research and Key Findings

OWN RESEARCH

Surveys

23

68% flagged hallucinated / inaccurate outputs

Student Interviews

5

Wrong citations went undetected until graded

Employee Interviews

3

Avoid AI for compliance-adjacent decisions

PUBLIC RESOURCES

- Adoption vs. real use: 86–91% “adopted” AI, only 10–15% use it daily
- Trust in AI accuracy stuck at 29–34% across all use cases
- Override rate: 74% under low transparency vs. 49% under high transparency
- #1 frustration everywhere: “almost right” outputs
- Regulated AI is most cautious — trust is earned per-tool, not by default

“I stopped trusting it after it gave me a citation that didn’t exist.”

— Student Interviewee, Round 1 Research

2

Hypothesis Tested

H1

Frustration is driven by hallucination, bias & accuracy issues — not speed.

✓ Validated — top complaint in surveys + secondary data

H2

Trust tracks transparency, not accuracy alone — users override what they can’t see the reasoning for.

✓ Validated — override rate 49% → 74% as transparency drops

H3

One risk model can’t fit all use cases — customer, internal & regulated AI need different oversight.

✓ Validated — trust ranges from 34% to 79% by context

H4

High adoption hides low real trust — usage looks solved when it isn’t.

✓ Validated — 86–91% adoption vs. only 10–15% daily use

3

The Insight

AI’s biggest failure isn’t a broken model — it’s an invisible one.

Across every context we studied — customer-facing, internal, and regulated — the same pattern repeats: broad adoption, shallow trust, and no real-time visibility into when AI goes wrong, not if.

This is a validated, structural trust gap — not a hypothetical one. ControlPlane.ai exists to close it: continuous, real-time oversight that makes AI’s blind spots visible and actionable across every use case.

RESEARCH BASE

23 surveys · 5 student interviews · 3 employee interviews · 5 secondary industry reports

COVERAGE

Spans all 3 enterprise AI contexts — customer-facing, internal, and regulated decision-support

AI adoption is fragmenting into three distinct populations
— oversight has to fit each one, not the loudest one



EXECUTIVE INSIGHT

ControlPlane.ai must work across three AI populations — each with its own inbuilt vs. standalone mix, user base, and reasons for reaching for AI. Getting this landscape right keeps oversight accurate for every use case.

WHY THIS MATTERS NOW



1.6% → 48%

Regulated AI usage growth



80% vs 34%

Trust: vetted vs. standalone AI



75% vs 51%

Leader vs. frontline adoption

1 Three AI populations organizations must govern — each with its own inbuilt vs. standalone mix



Customer-Facing AI

INBUILT AI

Company's own in-app assistant (bank/telco chatbot). Primary touchpoint for most end customers.

STANDALONE AI

General tools (ChatGPT, Gemini) used outside the company app — for research or a second opinion.

USER BASE (Standalone AI, by age)

18–34: 55% • 35–54: 35% • 55+: 10%

WHY THEY USE AI

- Instant, 24/7 answers — 82% prefer AI-first for immediate queries
- Convenience — 67% used a chatbot in the past year
- Still guarded — 79% prefer a human for serious issues



Internal Employee AI

INBUILT AI

Company-built copilot for employees only. Skews technical: Engineers 71%, Marketing 63%, IT 55%.

STANDALONE AI

ChatGPT/Copilot etc. used informally alongside sanctioned tools — often the default choice.

USER BASE (adoption gap)

Leaders: 75% vs. Frontline: 51% — the “silicon ceiling”

WHY THEY USE AI

- Saves time — 2.2 hrs/week per user reclaimed
- Offloads repetitive, low-value admin work
- Only 29% trust its accuracy (down from 40%) — 37% of time saved is lost fixing mistakes



Regulated Decision-Support

INBUILT AI

Vetted, compliance-approved tool embedded in the workflow. Highly preferred — 80% trust it.

STANDALONE AI

Third-party/general AI used ad hoc. Less preferred — discouraged for compliance-adjacent calls.

USER BASE (clinicians & compliance/finance)


48% of clinicians now use AI (up from 26%) — but just 16% for actual clinical decisions

WHY THEY USE AI

- 81% of physicians now use AI, 2x 2023's rate — mostly documentation & research
- Improves consistency & auditability of decisions
- Only 34% trust AI agent actions — override rate 74% (low transparency) vs. 49% (high)

BACKING THE LANDSCAPE — real growth, trust & preference data behind each use case

AI Usage Growth by Category (%)



Declining Trust in AI, 2023 → 2025 (%)



Regulated Workflows: Trust by AI Type (%)



2023 → 2025 (percentage-point change shown)



TAKEAWAY: One oversight model can't fit all three populations. Trust is falling for standalone use (customer & employee) while regulated AI use is growing fastest but least trusted for real decisions → governance must flex by population, tool type, and decision stakes.

ControlPlane.ai's risk taxonomy — and where each risk concentrates across the three AI populations


1 What can go wrong — the three risk dimensions ControlPlane Checker watches on every response

 PERFORMANCE <i>Confidently wrong responses</i>	 Prompt contains a mistaken / wrong established fact	 Missing important data in the prompt → hallucination	 No established resources to ground the answer → hallucinates
 RESPONSIBILITY <i>Bias · unsafe/illegal · leaking data</i>	 Bias: user-steered or personalization-driven skew	 Unsafe/illegal: harmful, toxic or confidential+illegal content	 Leaking data: private/org info surfaced from training data
 COST <i>What drives spend up</i>	 Multiple AI models used to cover each risk category	 Re-prompting flagged responses → infinite retry loops	 More verification layers → higher cost & slower AI

2 Where risk concentrates — highest and lowest exposure by AI population

 Customer-Facing AI 	 Internal Employee AI 	 Regulated Decision-Support 
 HIGHEST RISK: Bias Consumer-facing tone is steered by what the user wants to hear — personalization skew compounds over repeat sessions.	 HIGHEST RISK: Leaking Data Broad access to internal knowledge bases and confidential docs — easiest path for sensitive info to leave the org.	 HIGHEST RISK: Leaking Data Patient, financial & compliance records are the most sensitive data class — highest stakes if exposed.
ALSO ELEVATED: Hallucination High query volume on products/policies with thinggrounding — confidently wrong answers follow close behind bias.	ALSO ELEVATED: Cost Redundant use of inbuilt + standalone tools side-by-side inflates spend across a large employee base.	ALSO ELEVATED: Bias Fairness in lending, triage or compliance calls carries real consequences for real people.
 LOWEST RISK: Leaking Data Scoped mostly to the customer's own account/session — least exposure to deep organizational secrets.	 LOWEST RISK: Bias Mostly task-based (docs, code, admin) — less opinion-driven than consumer-facing conversation.	 LOWEST RISK: Hallucination Vetted, compliance-approved tools are tightly grounded — least prone to confidently-wrong answers.

⚡ TAKEAWAY: No population shares the same risk fingerprint — Customer-Facing skews toward bias/hallucination, Employee and Regulated AI both skew toward leaking data. ControlPlane Checker must weight its scoped, tiered checks differently per population, not apply one blanket rule.




TAKEAWAY: Both layers fork the same way — screen, classify, act — and every escalation or override is logged and fed back, sharpening the checkers over time.

ControlPlane.ai — core mechanism: a two-layer AI checker wrapped around every response*Prompt Analyzer screens intent before generation; Response Checker re-verifies after — every step shown to the user, nothing ships untracked***EXECUTIVE INSIGHT**

Every prompt runs through two independent AI layers that never trust each other, screening intent before the Core AI answers and verifying the output before a user ever sees it — with every check, tag and decision surfaced, not hidden in a black box.

WHY THIS MATTERS NOW **233 → 362**

documented AI incidents, 2024→2025 (Stanford HAI 2026 AI Index)

**1 Two layers, one shared design principle — screen before generation, verify after**

Prompt Analyzer (before): reads the whole prompt, checks every segment for unsafe intent, classifies safe segments into up to 6 categories, and resolves Wrong Prompt / Missing Data itself with an explicit instruction to Core AI.

Response Checker (after): re-screens the response for unsafe content, then checks only the tags that segment's category carries, and lands each response in Escalate, Edit or Reject — tags checked/crossed shown to the user.



TAKEAWAY: Screening before generation catches unsafe intent early and cheaply; verifying after generation catches what generation itself introduces — hallucination, bias, leakage. Neither layer alone closes the loop.

Detection technique follows the risk — six intent categories, scoped tag checks, then a graded decision

Layer 1 classifies and tags every safe segment; Layer 2 verifies only those tags and routes to Escalate, Edit or Reject

1 Layer 1 — Prompt Analyzer sorts every safe segment into up to 6 intent categories, each with its own risk tags



**Factual /
Verifiable**

Wrong Prompt · Missing Data



**Opinion /
Subjective**

Bias · Missing Data · Hallucination



Prediction

Hallucination · Bias



**Recommendation
/ Advice**

Bias



**Instruction /
Command**

Missing Data



Confidential

Leaking Data · Confidential+Illegal

Wrong Prompt: verified against known facts, passed to Core AI with a correction instruction — model answers informed by the real data, not the mistaken premise.

Missing Data: Core AI is instructed to do both — explicitly ask the user for the missing input, and still generate an answer built on stated assumptions.

2 Layer 2 — Response Checker verifies only the tags each segment carries, then routes to a graded decision



Hallucination

Checked against public sources + prompt's own data; unverifiable → explicit uncertainty note



Bias

Personalization / region / religion skew vs. prior context; confirmed → dual-window view



Leaking Data

Training-set or private info surfacing in the answer — any leak is a hard REJECT



ESCALATE

No confirmed risk — ships to user as generated



EDIT

Risk confirmed — annotated or dual-window split



REJECT

Universal fail / leak — hard stop, decline message



TAKEAWAY: REJECT is reserved for universal, non-negotiable failures — everything else is scoped. Tuning scrutiny to exactly the risk each segment could plausibly carry is what keeps false positives down and user trust survivable at scale.

Every decision leaves an audit trail — human review on ambiguous calls, feedback that retrains the system, live metrics on the home page

Nothing ships untracked: each override becomes tuning data for the next threshold

EXECUTIVE INSIGHT

Every EDIT-tier response, and every ambiguous confidential-data-plus-legality call, routes to optional human review. The reviewer's verdict — not a fixed rule — becomes tomorrow's confidence threshold, so the system ages with the risk landscape rather than against it.

WHY THIS MATTERS NOW

28% → 18%

orgs rating AI-incident response "excellent" (Stanford HAI / McKinsey, 2024→2025)

1 Audit trail — every segment, every tag checked, every verdict, reconstructable end-to-end

Segment	Category	Tags checked	Tags crossed	Verdict
"Will this dosage combo interact badly?"	Recommendation	Bias	Bias → confirmed	EDIT (dual-window)
"Draft summary using our Q3 dataset"	Confidential	Leaking Data	none crossed	ESCALATE

FEEDBACK LOOP

Every human override becomes tuning data — confidence thresholds and per-tag sensitivity adjust from real reviewer outcomes over time, not from a fixed rule that ages as models and regulation change.

2 Home-page metrics dashboard — trustworthiness in four numbers, updated live

48,206

prompts processed / wk

91%

user satisfaction

4.8%

false positive rate

1.3%

false negative rate



TAKEAWAY: Governance is not a separate slide from the checker — it's the same audit record, replayed as a dashboard: every allow, edit and block is logged, reviewable, and feeds the next threshold.

Sources: Stanford HAI 2026 AI Index Report · Times Voice of the SOC Analyst (71% of analysts report alert-fatigue burnout) — illustrative dashboard figures for prototype demonstration

ControlPlane.ai runs at the speed each use case needs — scoped checks, not a blanket tax on every response

Added latency is parallelized and tiered by risk tolerance; the highest-risk (unsafe/illegal) path is also the fastest

EXECUTIVE INSIGHT

The checker never runs one heavy pipeline per request. Segments are tagged once; only the tags a segment actually carries get verified — and Prompt Analyzer + Response Checker calls run concurrently, not sequentially. Real-time cases get a lightweight fast path; batch and human-reviewed workflows get deeper checks inside time they already budget.

WHY THIS MATTERS NOW

🕒 53% abandon >3s

of mobile users abandon a page/response taking longer than 3 seconds to return (Google/SOASTA). Latency tolerance — not just accuracy — is a real trust variable ControlPlane is designed around.

1 Baseline Core AI response time vs. time with ControlPlane in the loop



Customer-Facing AI

REAL-TIME · STRICT BUDGET

BASILINE	OVERHEAD	TOTAL
1.8s	+0.4s	2.2s

Fast-path scoped checks only (tags the segment's category actually carries) run in parallel with Response Checker — no sequential stacking.



Internal Employee AI

NEAR REAL-TIME

BASILINE	OVERHEAD	TOTAL
2.6s	+0.5s	3.1s

Same scoped-and-parallel design; slightly larger prompts (docs, code context) add a little Prompt Analyzer time.



Regulated Decision-Support

BATCH / HUMAN-PACED

BASILINE	OVERHEAD	TOTAL
Async	+0s	No add

Deepest verification runs inside the review window clinicians/analysts already take — checker time is absorbed, not added, not perceived.



The riskiest case is also the fastest: unsafe/illegal intent is caught by the universal check before segment classification even begins — it skips straight to REJECT, adding no extra latency at all.



TAKEAWAY: Overhead scales with risk tolerance, not the other way round — real-time surfaces get a thin, parallel scoped check (~20% add, sub-second), while batch and human-reviewed workflows get the deepest verification for effectively zero perceived cost.

Model-agnostic by design, tuned against alert fatigue, and reported the way a skeptical stakeholder can verify

Three design choices that keep ControlPlane trustworthy without ever touching model internals

1 Working at the input/output layer, overflagging discipline, and stakeholder-grade metrics



Input/Output layer, not model internals

Why it's the only option: Enterprises consume a foundation model via API — GPT, Gemini, Claude, or an in-house model. There is no access to weights, attention maps, or logits to inspect.

What we check instead: The prompt going in and the response coming out — the two points every enterprise actually controls, regardless of which model sits underneath.

The trade-off, named honestly: We can't cite an internal confidence score. We offset this with multiple independent external signals — segment tagging, tag-specific verification, retrieval grounding, human review — instead of one internal one.

Why this is a strength, not a workaround: Swap Core AI for a different model tomorrow and the checker needs zero re-engineering — the boundary it watches never moves.



Overflagging — engineered out, not just avoided

4.8% false positive rate today — tuned, not accidental

How we keep it low: Response Checker verifies only the tags each segment's category actually carries (e.g. only Bias for a Recommendation) — not every dimension on every response — so scrutiny matches plausible risk instead of over-checking by default.

Graded, not binary: Escalate / Edit / Reject means most confirmed risk is annotated inline (EDIT), not blocked — preserving flow instead of a hard stop that trains users to route around us.

If overflagging did creep in: Users start overriding or ignoring warnings; the override-rate research already shows this pattern (74% under low transparency vs. 49% high) — the very trust gap ControlPlane exists to close would reopen.

Self-correcting over time: Every reviewer override retrains classifier thresholds — a tag that fires too often gets tightened automatically, not just noted.



Metrics that satisfy a skeptical stakeholder

What's defined and measured: False positive = flagged, but human reviewer approved as-is. False negative = shipped, but later found (via retrospective audit sampling) to carry hallucination, bias or a leak.

How it's reported: Live ops dashboard (FP/FN rate, override rate, volume, satisfaction) for engineers; a weekly trend summary for business stakeholders; a quarterly drift review asking whether thresholds still fit as models and regulation change.

The pitch to a skeptic: Not "trust us" — "audit us." Every verdict is reconstructable end-to-end: segment → tags checked → tags crossed → verdict, the same record the dashboard is built from.

Why that's the stronger claim: A single accuracy number can be cherry-picked; a reconstructable audit trail lets the stakeholder check our claim themselves, on any case they choose.




TAKEAWAY: None of this depends on seeing inside the model — it depends on disciplined design at the boundary: scoped checks that don't over-flag, and an audit trail transparent enough for a skeptic to verify rather than take our word for it.


From working prototype to enterprise rollout — a phased implementation roadmap

The core mechanism is already built and demonstrated; what remains is sequencing adoption, not re-engineering

1 Implementation roadmap — four phases from pilot to enterprise-wide governance



2 Prioritized recommendations — what to do first, next, and later



Sequencing beats scope: prove the mechanism on one high-risk workflow, let real overrides tune the thresholds, then extend the same core — unchanged — to every other use case.

Illustrative phased rollout for prototype demonstration

Built to prove itself — expected outcomes, managed risk, and a defensible edge

The same audit data that powers the dashboard also sets the targets, names the risks, and proves the differentiation

1 KPI & expected outcomes — from prototype baseline to production targets


4.8% → <3%

False positive rate
fewer unnecessary escalations


1.3% → <0.5%

False negative rate
near-zero undetected risk


-30%

Human review load
feedback-tuned thresholds


100%

Pre-response coverage
regulated flows checked before ship

2 Risk management & why ControlPlane.ai wins

RISK

MITIGATION — ALREADY IN THE PIPELINE

Over-flagging drives alert fatigue

Tiered Escalate / Edit / Reject, tuned by the feedback loop — not static rules

No reliable real-time ground truth to verify claims

Retrieval verification against source docs + explicit uncertainty tagging

Regulation varies by geography / industry and keeps evolving

Configurable, versioned policy layer per use case & geography, fully audit-logged

Latency budget breach when scaling to real-time channels

Scoped, parallelized checks — only the tags a segment carries; async post-hoc audit elsewhere



WHY CONTROLPLANE.AI

vs. point solutions: one two-layer pipeline covers bias, hallucination and privacy — not three stitched-together vendors

vs. static guardrails: thresholds retrain from real reviewer overrides, so rules age with the model and regulation, not against them

vs. black-box filters: works at the input/output layer — model-agnostic for enterprises that consume AI via API

vs. manual-only review: scales to tens of thousands of interactions/week; humans touch only ambiguous EDIT-tier cases



The prototype is the mitigation, not a promise: every risk above already has a named control engineered into the pipeline demonstrated in Slides 1–3.

Thank you

